

TECHNICAL RELEASE

Segovia v0.4.1: a bounded-memory, near-real-time streaming preprocessor for Neuropixels-scale electrophysiology

Felipe Carvajal Brown^{1,*}¹Independent Researcher

*fcarvajalbrown@gmail.com; ORCID 0000-0002-8300-7587

Abstract

Segovia is an open-source Rust library with Python bindings for bounded-memory, near-real-time preprocessing of high-density extracellular electrophysiology recordings. It applies a bandpass-filter, common-median-reference, and whitening chain to Neuropixels-scale (384–768 channels, 30 kHz) 16-bit integer streams one chunk at a time, releasing the Python Global Interpreter Lock for each Rust computation. Peak memory is bounded by chunk size, not by recording length, and scales analytically. On a full 55.8-minute International Brain Laboratory AP-band recording (385 channels, mtscomp-compressed) at a 300 ms chunk budget, Segovia holds 99.7% real-time deadline adherence at 0.21 GB peak resident memory. The package ships three production readers (SpikeGLX, Zarr, mtscomp) and a built-in streaming synthetic simulator.

Availability and Implementation: Source code (Rust and Python, AGPL-3.0-or-later):

<https://github.com/fcarvajalbrown/Segovia>. Install via `pip install segovia` (PyPI) or `cargo add segovia` (crates.io). Version v0.4.1.

Key words: electrophysiology; Neuropixels; streaming; real-time preprocessing; bounded memory; Rust; spike sorting

Statement of Need

Neuropixels probes [1] acquire 384 channels at 30 kHz (~22 MB/s per probe). Standard preprocessing pipelines (SpikeInterface [2], MountainSort) run offline on completed recordings and are optimized for batch throughput. Near-real-time applications, such as closed-loop stimulation and online brain-machine interface decoding, require that each incoming chunk of samples be preprocessed within its acquisition period (the real-time deadline), with peak memory independent of recording length. Batch tools driven one chunk at a time re-read filter-margin neighbourhoods on every call and do not bound memory analytically. Segovia was built to fill this gap: a composable streaming preprocessor whose memory ceiling is fixed by chunk size and whose per-chunk latency meets the acquisition deadline on real recordings.

Implementation

Table 1. Real IBL AP-band recording, full length (385 channels, mtscomp-compressed, 55.8 min, 11,167 chunks, 300 ms budget, steady state).

Engine	Mean	p99	Max	Adherence	Peak RSS	Jitter
Segovia	179.2 ms	277.0 ms	334.5 ms	99.7%	0.21 GB	38.6 ms
SpikeInterface online	205.3 ms	355.0 ms	932.0 ms	94.7%	0.41 GB	60.5 ms

Table 2. Real IBL AP-band recording, cold-start first 60 s (385 channels, mtscomp-compressed).

Chunk	Engine	Mean latency	p99	Adherence	Peak RSS
100 ms	Segovia	92.9 ms	122.0 ms	74.2%	0.21 GB
100 ms	SpikeInterface online	112.0 ms	275.2 ms	64.2%	0.46 GB
300 ms	Segovia	194.5 ms	256.4 ms	100%	0.28 GB
300 ms	SpikeInterface online	245.8 ms	365.7 ms	69.5%	0.52 GB
1000 ms	Segovia	617.3 ms	705.9 ms	100%	0.49 GB
1000 ms	SpikeInterface online	786.0 ms	947.5 ms	98.2%	0.74 GB

Core abstraction

Segovia’s `ChunkSource` trait is an iterator over fixed-size (samples $\times$ channels) 16-bit integer buffers. Three production implementations are provided:

- `SpikeGlxReader`: memory-mapped SpikeGLX `.bin + .meta` (zero-copy).
- `ZarrReader`: chunked Zarr arrays (gzip, zstd, Blosc) via the `zarrs` crate.
- `CbinReader`: mtscomp-compressed IBL `.cbin`, per-chunk zlib decompression via `flate2`.

Preprocessing chain

`preprocess(chunk_source, config)` applies a Rayon-parallelized chain: a 5th-order Butterworth bandpass filter (zero-phase `sosfiltfilt`), common median reference, and global ZCA whitening. The Python GIL is released via `PyO3`’s `allow_threads` for the Rust computation. Cross-chunk IIR filter state is maintained deterministically regardless of thread count.

Peak resident memory scales as $\text{batch} \times (\text{chunk} + 2 \text{ margin}) \times \text{channels} \times \text{sizeof}(f32)$ and is independent of recording length. The default (auto) batch width is capped at $\min(\text{logical threads}, 4)$ as a conservative out-of-memory guard, bounding the default footprint to ~ 1.2 GB on any machine; callers may pin an explicit batch. This cap is an out-of-memory safety guard, not a throughput-optimum claim.

Built-in simulator

`SyntheticEphysReader` emits arbitrarily long synthetic streams without writing to disk. Spike templates use a Ricker temporal waveform with point-source spatial decay ($V(r) = A d_{\perp} / r$); firing is Poisson per unit; noise is additive white Gaussian. The pseudorandom number generator (SplitMix64 + xoshiro256++) is platform-independent and chunk-size-independent. `ground_truth()` returns (sample, unit, channel) tuples for downstream sorter evaluation.

Results

Evaluation follows a replay-at-acquisition-rate methodology: data are streamed at the true 30 kHz rate with per-chunk compute latency measured by a monotonic clock. Deadline adherence is the fraction of chunks with latency at or below the chunk period.

Table 1 reports the full-length steady-state comparison. At steady state Segovia leads on every axis; the decisive margins are peak memory ($2\times$), maximum latency ($2.8\times$), and jitter. Table 2 is the cold-start first-60 s window, where SpikeInterface’s warm-up cost is highest and the adherence gap is widest (100% vs 69.5% at 300 ms); that gap narrows to the steady-state figures over the full recording.

On synthetic recordings (384 channels, `SyntheticEphysReader`, seed 0) Segovia reaches 100% deadline adherence at all chunk sizes (100/300/1000 ms) with jitter 3.6/8.6/18.6 ms.

Peak memory is bounded and file-size-independent on both sources; on the full 55.8-minute (29 GB compressed) real recording the memory bound holds to within 1% of a ten-minute slice. In a batch-throughput comparison on that full recording, Segovia at a pinned batch held peak memory below both SpikeInterface executor modes (1.19 GB vs 2.18 GB thread-pool and 4.42 GB process-pool) and completed in less wall time (806 s vs 923 s and 1022 s) in a single run; the memory bound is the robust result and the wall-time margin warrants multi-run replication. Throughput exceeds the 22 MB/s NeuroPixels acquisition rate at all configurations.

Known limitation: the zero-phase Butterworth filter introduces a 50 ms look-ahead; a causal filter mode is not yet implemented. Benchmarks are on a single machine (Windows, 8 physical / 16 logical cores).

Availability of source code and requirements

- Project name: Segovia
- Project home page: <https://github.com/fcarvajalbrown/Segovia>
- Operating system(s): Linux, macOS, Windows (platform independent)

- Programming language: Rust (core, edition 2021, minimum supported Rust version 1.74); Python bindings via PyO3
- Other requirements: Python ≥ 3.8 for the bindings (prebuilt abi3 wheels on PyPI); a Rust toolchain and `maturin` to build from source
- License: AGPL-3.0-or-later (OSI-approved)
- RRID: to be registered at SciCrunch.org prior to publication
- bio.tools ID: not yet registered
- Version: v0.4.1

The source code is under the AGPL-3.0-or-later license (an Open Source Initiative approved license) and is openly available in the GitHub repository above, with versioned releases on PyPI and crates.io.

Availability of Test Data

Segovia is a software library; no new experimental dataset was generated. The real recording used for evaluation is a publicly released International Brain Laboratory AP-band Neuropixels recording (`_spikeglx_ephysData_g0_t0.imec0.ap.cbin`), obtainable from the International Brain Laboratory data portal. Synthetic recordings are generated deterministically on demand by `SyntheticEphysReader` (seed-fixed, no deposit required). The form of the required GigaDB deposit for a software Technical Release is being confirmed with the editorial team.

Declarations

List of abbreviations

AGPL: GNU Affero General Public License; AP: action potential (band); CMR: common median reference; GIL: Global Interpreter Lock; IBL: International Brain Laboratory; IIR: infinite impulse response; RSS: resident set size; SOS: second-order sections; ZCA: zero-phase component analysis (whitening).

Ethical Approval

Not applicable. No human or animal data were collected for this work; the recording used for evaluation is a publicly released dataset.

Competing Interests

The author declares that they have no competing interests.

Funding

No funding was received for this work.

Author's Contributions

E.C.B. is the sole author and performed all work (CRediT: conceptualization, software, methodology, investigation, writing).

Acknowledgements

The author thanks the International Brain Laboratory for the publicly available recording used in evaluation. The software is named in memory of Claudio Segovia.

References

1. Jun JJ, Steinmetz NA, Siegle JH, Denman DJ, Bauza M, Barbarits B, et al. Fully integrated silicon probes for high-density recording of neural activity. *Nature* 2017;551(7679):232–236.
2. Buccino AP, Hurwitz CL, Garcia S, Magland J, Siegle JH, Hurwitz R, et al. SpikeInterface, a unified framework for spike sorting. *eLife* 2020;9:e61834.